



Spring Boot Microservices

Beginner to Guru

Distributed Tracing



What is Distributed Tracing?

- Monoliths have the luxury of being self contained, thus tracing typically is not needed
- Transactions in microservices can span many services / instances, and even data centers
- Distributed tracing provides the tools to trace a transaction across services and nodes
- Distributed tracing is used for two aspects:
 - Performance monitoring across steps
 - Logging / troubleshooting





Spring Cloud Sleuth

- Spring Cloud Sleuth is the distributed tracing tool for Spring Cloud
- Spring Cloud Sleuth uses an open source distributed tracing library called Brave
- Conceptually what happens:
 - Headers on HTTP requests or messages are enhanced with trace data
 - Logging is enhanced with trace data
 - Optionally trace data can be reported to Zipkin

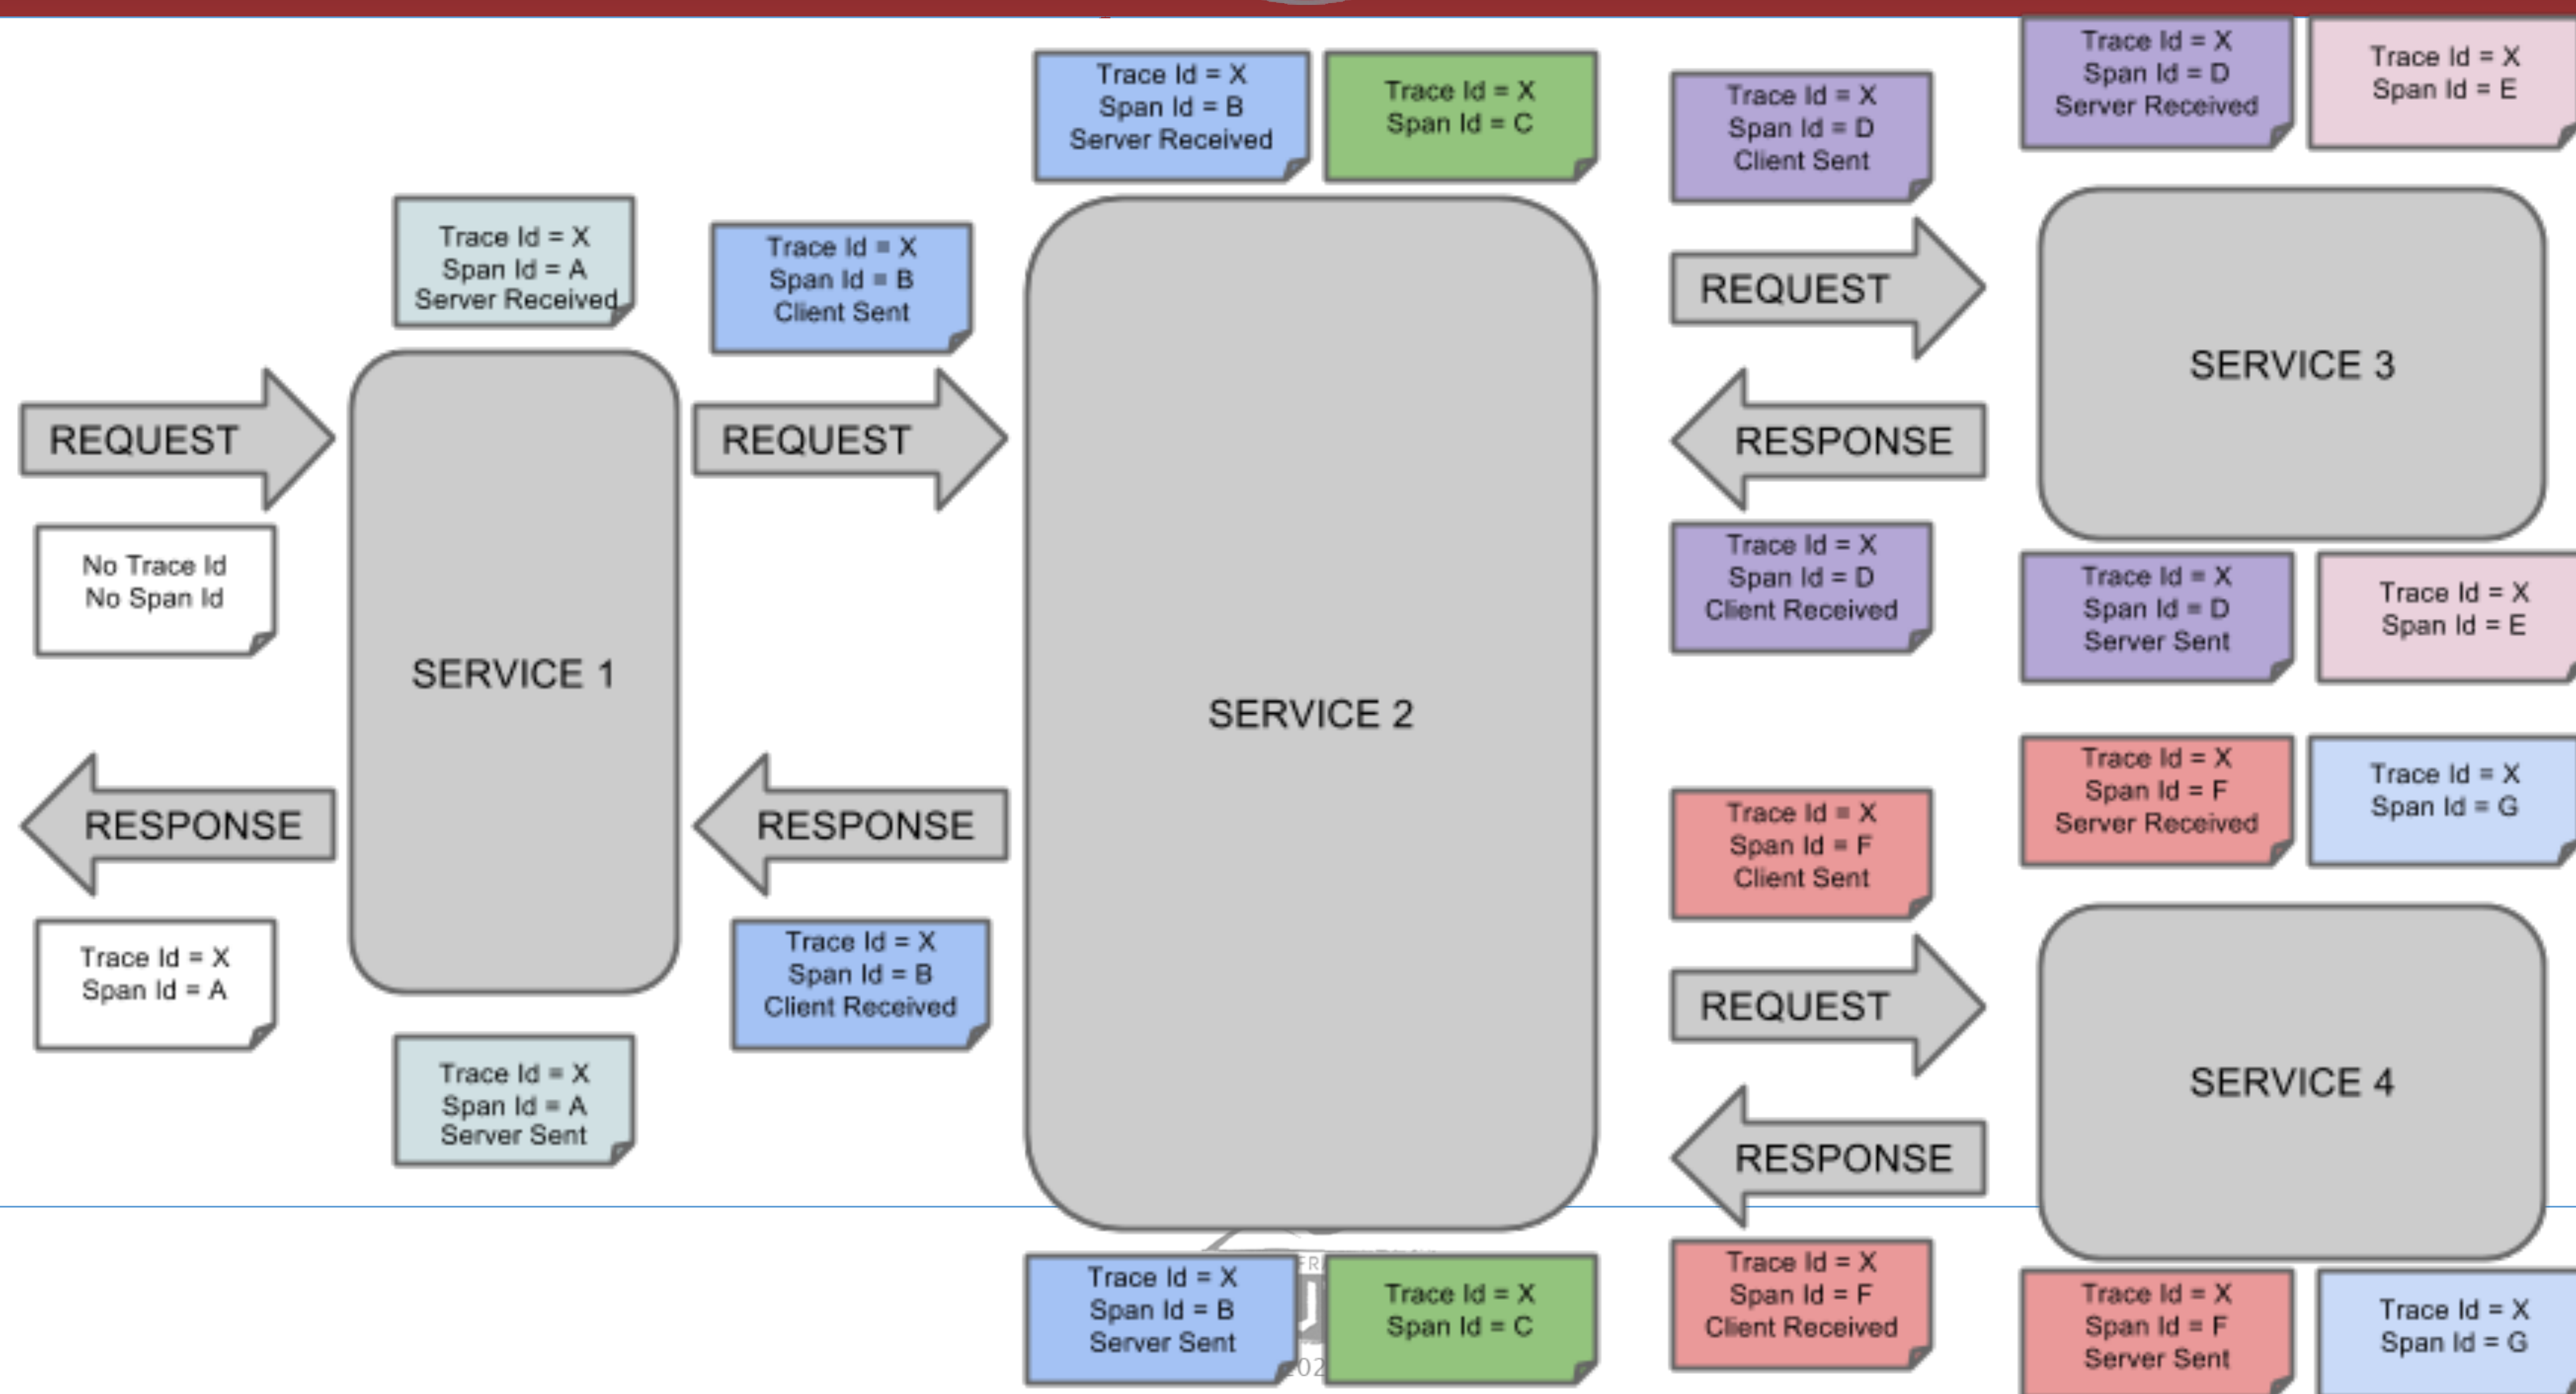


Tracing Terminology

- Spring Cloud Sleuth uses terminology established by Dapper
 - Dapper is a distributed tracing system created by Google for their production systems
- **Span** - is a basic unit of work. Typically a send and receive of a message.
- **Trace** - A set of spans for a transaction
- **cs / sr** - Client Sent / Server Received - aka the request
- **ss / cr** - Server Sent / Client Received - aka the response

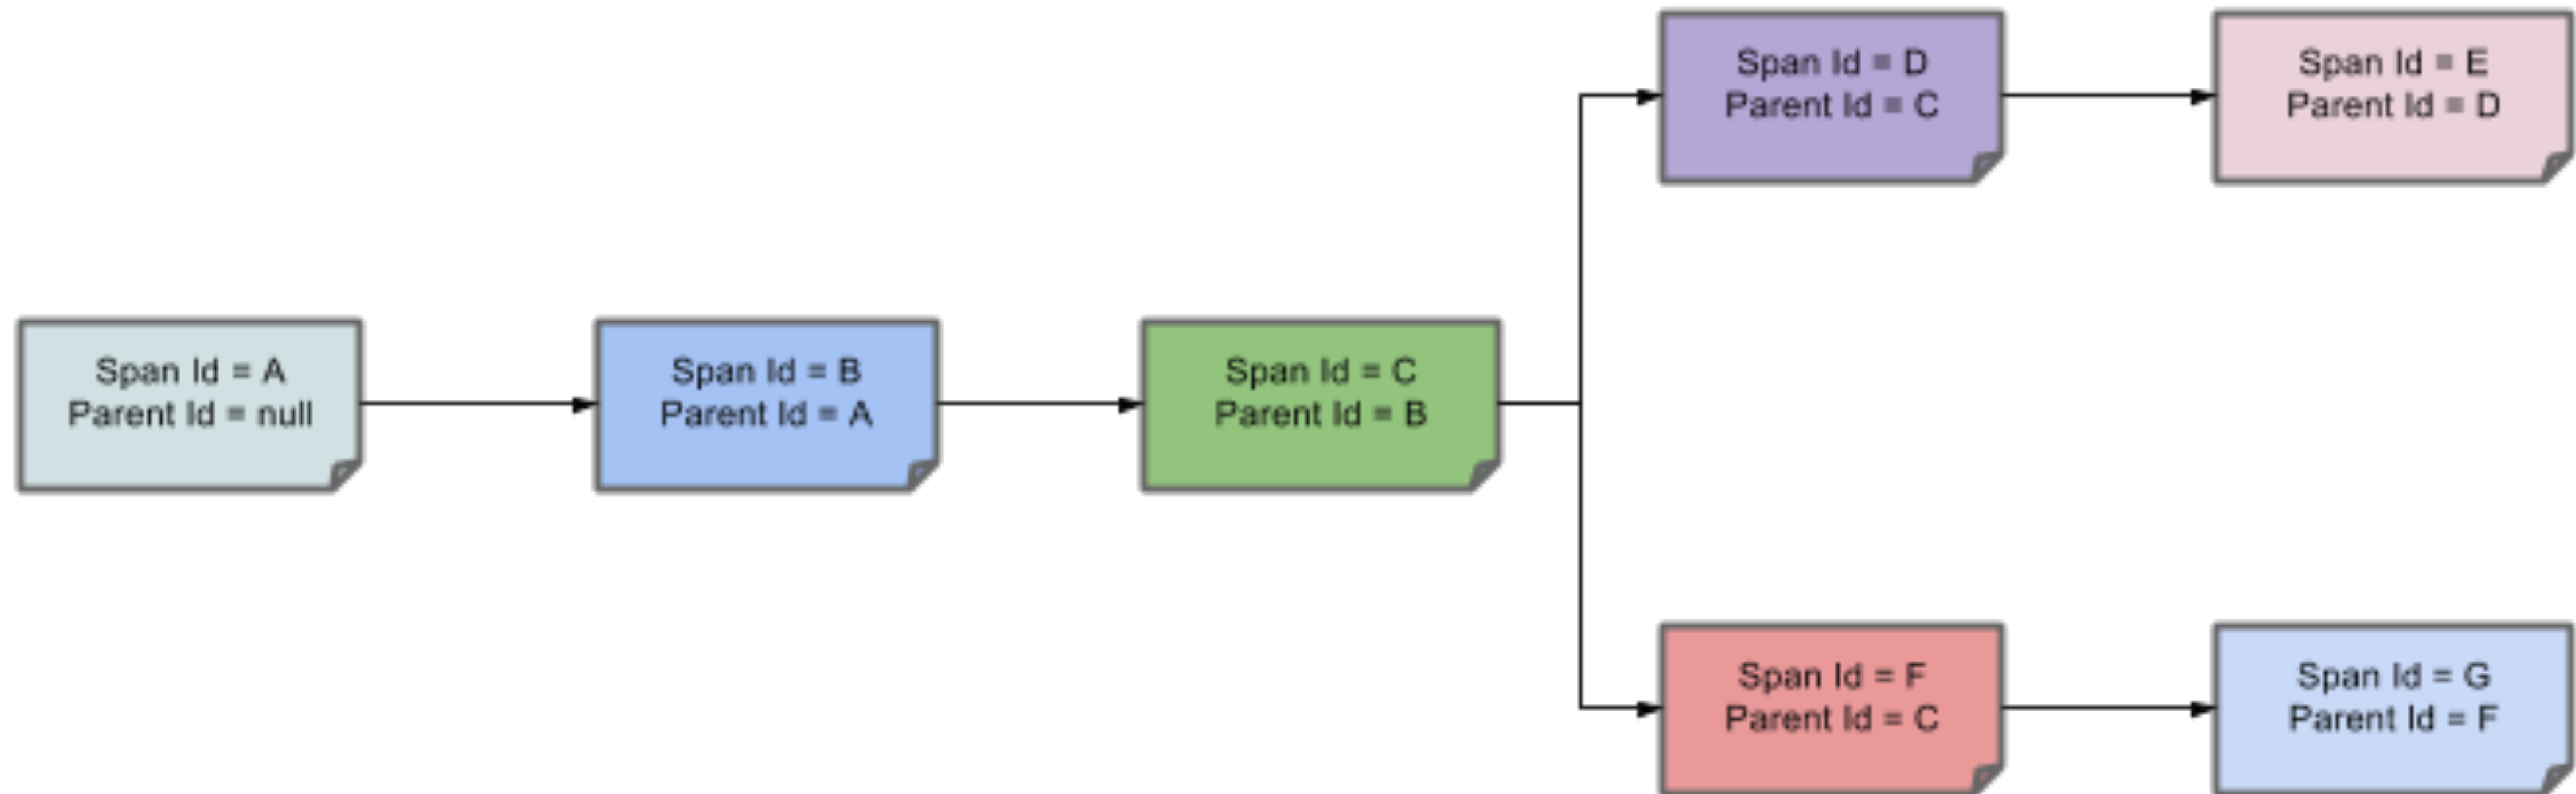


Source Pivotal





Source Pivotal





Zipkin Server

- Zipkin is an open source project used to report distributed tracing metrics
- Information can be reported to Zipkin via webservices via HTTP
 - Optionally metrics can be provided via Kafka or Rabbit
- Zipkin is a Spring MVC project
 - Recommended to use binary distribution or Docker image
 - Building your own is not supported
- Uses in memory database for development
 - Cassandra or Elasticsearch should be used for production





Zipkin Quick Start

- Via Curl:
 - `curl -sSL https://zipkin.io/quickstart.sh | bash -s`
 - `java -jar zipkin.jar`
- Via Docker (Recommend for course):
 - `docker run -d -p 9411:9411 openzipkin/zipkin`
- View traces in UI at:
 - `http://your_host:9411/zipkin/`





Installing Spring Cloud Sleuth

- `org.springframework.cloud:spring-cloud-starter-sleuth`
 - Starter for logging only
- `org.springframework.cloud:spring-cloud-starter-zipkin`
 - Starter for Sleuth with Zipkin - includes Sleuth dependencies
- Property `spring.zipkin.baseUrl` is used to configure Zipkin server



Logging Output

- **Example:** - `DEBUG [beer-service,39853b63c1c3f919,419b9ac9a073bbba,true]`
 - [Appname, Traceld, SpanId, exportable]
- **Appname** - Spring Boot Application Name
- **Traceld** - Id value of the trace
- **SpanId** - Id of the Span
- **Exportable** - Should span be exported to Zipkin? (Programmatic configuration option)



Logging Configuration

- Microservices typically will use consolidated logging
- Number of different approaches for this - highly dependent on deployment environment
- Consolidated logging will be covered in a future section of the course
- To support consolidated logging, log data should be available in JSON
- Spring Boot by default uses logback, which is easy to configure for JSON output
 - To be covered in separate lesson

