



Spring Boot Microservices

Beginner to Guru


The Need for Sagas




The Microservice Death Star

- As the number of services grows, so does complexity
- Death Star - Examples
 - Netflix
 - Twitter
 - Uber









Challenges

- Business transactions often span multiple microservices
- ACID transactions are not an option between services
- Distributed Transactions / Two Phase Commits
 - Complex and do not scale
- Microservices should be technology agnostic
 - Making 2PC even more difficult to implement




CAP Theorem

- **CAP** - Consistency, Availability, Partition Tolerance
 - **Consistency** - Every read will have the most recent write
 - **Availability** - Each read will get a response, but without the guarantee data is most recent write
 - **Partition Tolerance** - System continues in lieu of communications errors or delays
- **CAP Theorem** - States a distributed system can only maintain two of three
 - Consistency, Availability, Partition Tolerance



CA Category
RDBMS



CP Category
BigTable
HBase
MongoDB
Redis

AP Category
Dynamo
Voldemort
Cassandra
CouchDB



2019



BASE - An ACID Alternative

- **BASE** - Coined by Dan Pritchett of EBay in 2008
- **BASE** = **B**asically **A**vailable, **S**oft state, **E**ventually consistent
 - The opposite of ACID
- **Basically Available** - Build system to support partial failures
 - Lost of some functionality vs total system loss.
- **Soft State** - Transactions cascade across nodes, it can be inconsistent for a period of time
- **Eventually Consistent** - When processing is complete, system will be consistent



Feral Concurrency Control

- *Feral Concurrency Control: An Empirical Investigation of Modern Application Integrity*
 - Published by Peter Baillis in 2015
- **Feral Concurrency Control** - Are application level mechanisms for maintaining database integrity
 - Relational Databases can enforce a variety of constraints - such as foreign key constraints
 - Not available within a distributed system
 - Thus up to the application to enforce constraints



Introducing Sagas

- Concept introduced in 1987 by Gracia-Molina / Salem of Princeton University
- Originally was looking at Long Lived Transactions (LLTs) within a single database
 - LLTs hold on to database resources for an extended period of time
 - In 1987 computational resources were much more scarce
- Paper proposed rather than long complex processes to break up into smaller more atomic transactions
- Introduced the concept of compensating transactions to correct partial executions



Sagas

- Sagas are simply a series of steps to complete a business process
- Sagas coordinate the invocation of microservices via messages or requests
- Sagas become the transactional model
- Each step of the Saga can be considered a request
- Every step of the Saga has a compensating transaction (request)
 - Semantically undoes the effect of the request
 - Might not restore to the exact previous state - but effectively the same



Saga Steps

- Each step should be a message or event to be consumed by a microservice
- Steps are asynchronous
- Within a microservice, it's normal to use traditional database transactions
- Each message (request) should be idempotent
 - Meaning if same message / event is sent there is no adverse effect on system state
- Each step has a compensating transaction to undo the actions



Compensating Transactions

- Compensating Transactions
 - Effectively become the 'Feral Concurrency Control'
 - Are the mechanism used to maintain system integrity
 - Should also be Idempotent
 - Cannot abort - need to ensure proper execution
 - Not the same as a 'rollback' to the exact previous state
 - Implements business logic for a failure event



Sagas are ACD

- **A - Atomic**
 - All transactions are executed or compensated
- **C - Consistency**
 - Referential integrity within a service by the local database
 - Referential integrity across services by the application - 'Feral Concurrency Control'
- **D - Durability**
 - Persisted by database of each microservice



Sagas & Eventually Consistent

- **BASE** = **B**asically **A**vailable, **S**oft state, **E**ventually consistent
- During execution of the Saga system is in the 'Soft State'
- Eventually consistent - meaning system will be consistent at the conclusion of the Saga
 - Consistency achieved via normal completion of the Saga
 - In the event of an error, consistency is achieved via compensating transactions



Summary

- **Saga Definition** - *"A long, involved story, account, or series of incidents"*
- Microservices by nature run in a distributed environment, across many computers
- Sagas are used to address the challenges faced when operating in a distributed environment
- Sagas are a tool used to coordinate a series of steps for a business transaction across multiple services
- Sagas not only prescribe the series of necessary steps, they also maintain system integrity

